

Assignment 9, Question 4, Amogha Sekhar, A53301791, CSE 250A

```
In [1]: import numpy as np
        from collections import defaultdict
```

```
In [2]: def read_file(filename):
        array_prob = np.zeros((81, 81), dtype = np.float)
        with open(filename, 'r') as f:
            for line in f:
                s, s_, prob = line.strip().split() # s_ is s'
                s, s_, prob = int(s)-1, int(s_)-1, float(prob)
                array_prob[s, s_] = prob
        return array_prob
```

```
In [3]: prob_a1 = read_file("prob_a1.txt")
        prob_a2 = read_file("prob_a2.txt")
        prob_a3 = read_file("prob_a3.txt")
        prob_a4 = read_file("prob_a4.txt")
```

```
In [4]: reward = np.zeros((81, 1), dtype = np.float)
        with open('rewards.txt', 'r') as f:
            i = 0
            for line in f:
                reward[i] = float(line.strip())
                i += 1
```

```
In [5]: gamma = 0.9925
```

Policy Iteration

```

In [6]: policy = np.zeros((81, 1), dtype = np.int) # The policy initialized at random
v_pi = np.ones((81, 1), dtype = np.float) # Value function
v_pi_before = np.zeros((81, 1), dtype = np.float) # Value function of prior

iter_ = 0

while max(abs(v_pi - v_pi_before)) > 0.0001:

    p_pi = np.zeros((81, 81), dtype = np.float)

    for s in range(81):
        for s_ in range(81):
            if policy[s] == 0:
                p_pi[s, s_] = prob_a1[s, s_]
            if policy[s] == 1:
                p_pi[s, s_] = prob_a2[s, s_]
            if policy[s] == 2:
                p_pi[s, s_] = prob_a3[s, s_]
            if policy[s] == 3:
                p_pi[s, s_] = prob_a4[s, s_]

    v_pi_before = np.copy(v_pi)
    v_pi = np.linalg.inv(np.identity(81) - gamma * p_pi).dot(reward)

    # argmax part
    for s in range(81):

        sum_a1, sum_a2, sum_a3, sum_a4 = 0, 0, 0, 0

        for s_ in range(81):
            sum_a1 += prob_a1[s, s_] * v_pi[s_]
            sum_a2 += prob_a2[s, s_] * v_pi[s_]
            sum_a3 += prob_a3[s, s_] * v_pi[s_]
            sum_a4 += prob_a4[s, s_] * v_pi[s_]

        list_sum = [sum_a1, sum_a2, sum_a3, sum_a4]
        best_policy = list_sum.index(max(list_sum))

        policy[s] = best_policy

        iter_ += 1

print("Took", iter_, "iterations to converge")

```

Took 486 iterations to converge

```
In [7]: v_pi
```

```
Out[7]: array([[ 0.          ],
 [ 0.          ],
 [100.70098073],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [102.3752644 ],
 [101.52364515],
 [ 0.          ],
 [ 0.          ],
 [109.48993454],
 [110.40903296],
 [111.33584663],
 [ 0.          ],
 [ 0.          ],
 [103.23462342],
 [ 0.          ],
 [106.77826755],
 [107.67462643],
 [108.57848712],
 [ 0.          ],
 [112.27044032],
 [ 0.          ],
 [ 0.          ],
 [104.10121204],
 [104.97507555],
 [105.88853591],
 [ 0.          ],
 [ 0.          ],
 [114.1632295 ],
 [113.21287932],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [103.78140737],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [115.12155727],
 [ 0.          ],
 [ 0.          ],
 [ 0.          ],
 [-133.33333333],
 [ 90.9853796 ],
 [-133.33333333],
 [ 0.          ],
 [-133.33333333],
 [116.08792959],
 [122.02491241],
 [ 0.          ],
 [ 0.          ]]
```

```
[ 81.39949278],  
[ 93.67165583],  
[ 95.17285726],  
[108.34261934],  
[109.58365072],  
[123.64307021],  
[123.1822391 ],  
[ 0.          ],  
[ 0.          ],  
[-133.33333333],  
[ 81.39949278],  
[-133.33333333],  
[ 0.          ],  
[-133.33333333],  
[125.24978944],  
[124.20738563],  
[ 0.          ],  
[ 0.          ],  
[ 0.          ],  
[ 0.          ],  
[ 0.          ],  
[ 0.          ],  
[ 0.          ],  
[ 0.          ],  
[133.33333333],  
[ 0.          ],  
[ 0.          ]])
```

```
In [9]: policy
```

```
Out[9]: array([[0],
               [0],
               [2],
               [0],
               [0],
               [0],
               [0],
               [0],
               [0],
               [0],
               [2],
               [1],
               [0],
               [0],
               [3],
               [3],
               [2],
               [0],
               [0],
               [2],
               [0],
               [3],
               [3],
               [0],
               [0],
               [2],
               [0],
               [0],
               [3],
               [3],
               [0],
               [0],
               [0],
               [2],
               [1],
               [0],
               [0],
               [0],
               [0],
               [0],
               [0],
               [0],
               [2],
               [0],
               [0],
               [0],
               [0],
               [0],
               [0],
               [0],
               [0],
               [2],
               [2],
               [0],
               [0],
```

```
[3],  
[3],  
[3],  
[3],  
[3],  
[2],  
[2],  
[0],  
[0],  
[0],  
[0],  
[0],  
[0],  
[0],  
[0],  
[2],  
[1],  
[0],  
[0],  
[0],  
[0],  
[0],  
[0],  
[0],  
[0],  
[0],  
[0]])
```

Value Iteration

```
In [10]: v_pi = np.ones((81, 1), dtype = np.float) # Value function
v_pi_before = np.zeros((81, 1), dtype = np.float) # Value function of prior

iter_ = 0

while max(abs(v_pi - v_pi_before)) > 0.000001:

    for s in range(81):

        sum_a1, sum_a2, sum_a3, sum_a4 = 0, 0, 0, 0

        for s_ in range(81):
            sum_a1 += prob_a1[s, s_] * v_pi_before[s_]
            sum_a2 += prob_a2[s, s_] * v_pi_before[s_]
            sum_a3 += prob_a3[s, s_] * v_pi_before[s_]
            sum_a4 += prob_a4[s, s_] * v_pi_before[s_]

        v_pi_before[s] = v_pi[s]

        list_sum = [sum_a1, sum_a2, sum_a3, sum_a4]
        v_pi[s] = reward[s] + gamma * max(list_sum)

    iter_ += 1

print("Took", iter_, "iterations to converge")
```

Took 3672 iterations to converge

```
In [11]: v_pi
```

```
Out[11]: array([[ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 1.00700865e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 1.02375148e+02],
 [ 1.01523530e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 1.09489817e+02],
 [ 1.10408916e+02],
 [ 1.11335729e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 1.03234507e+02],
 [ 9.93633760e-07],
 [ 1.06778151e+02],
 [ 1.07674510e+02],
 [ 1.08578371e+02],
 [ 9.93633760e-07],
 [ 1.12270323e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 1.04101096e+02],
 [ 1.04974959e+02],
 [ 1.05888420e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 1.14163112e+02],
 [ 1.13212762e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 1.03781292e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 1.15121440e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
 [-1.33333200e+02],
 [ 9.09852775e+01],
 [-1.33333200e+02],
 [ 9.93633760e-07],
 [-1.33333200e+02],
 [ 1.16087812e+02],
 [ 1.22024788e+02],
 [ 9.93633760e-07],
 [ 9.93633760e-07],
```



```
[ 8.13994071e+01],
[ 9.36715583e+01],
[ 9.51727592e+01],
[ 1.08342509e+02],
[ 1.09583540e+02],
[ 1.23642946e+02],
[ 1.23182115e+02],
[ 9.93633760e-07],
[ 9.93633760e-07],
[-1.33333200e+02],
[ 8.13994072e+01],
[-1.33333200e+02],
[ 9.93633760e-07],
[-1.33333200e+02],
[ 1.25249665e+02],
[ 1.24207261e+02],
[ 9.93633760e-07],
[ 9.93633760e-07],
[ 9.93633760e-07],
[ 9.93633760e-07],
[ 9.93633760e-07],
[ 9.93633760e-07],
[ 9.93633760e-07],
[ 9.93633760e-07],
[ 1.33333202e+02],
[ 9.93633760e-07],
[ 9.93633760e-07]])
```

```
In [12]: policy = np.zeros((81, 1), dtype = np.int) # The policy initialized at rand

for s in range(81):
    sum_a1, sum_a2, sum_a3, sum_a4 = 0, 0, 0, 0

    for s_ in range(81):
        sum_a1 += prob_a1[s, s_] * v_pi[s_]
        sum_a2 += prob_a2[s, s_] * v_pi[s_]
        sum_a3 += prob_a3[s, s_] * v_pi[s_]
        sum_a4 += prob_a4[s, s_] * v_pi[s_]

    list_sum = [sum_a1, sum_a2, sum_a3, sum_a4]
    best_policy = list_sum.index(max(list_sum))

    policy[s] = best_policy
```

```
In [13]: policy
```

```
Out[13]: array([[0],  
                [0],  
                [2],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [2],  
                [1],  
                [0],  
                [0],  
                [3],  
                [3],  
                [2],  
                [0],  
                [0],  
                [2],  
                [0],  
                [3],  
                [3],  
                [0],  
                [0],  
                [2],  
                [0],  
                [0],  
                [3],  
                [3],  
                [0],  
                [0],  
                [0],  
                [2],  
                [1],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [2],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [0],  
                [2],  
                [2],  
                [0],  
                [0],
```

In []: